

ECE 759
High Performance Computing for Engineering Applications
Assignment 6
Due Friday 4/14/2025 at 23:59 PM

Submit all plots on Canvas. Do not zip your Canvas submission. All *source files* should be stored in the **HW06** folder on the **main** branch of your **GitHub** repo. Please use the name **HW06** exactly as shown here (both in terms of capitalization & name). The **HW06** subdirectory should have no subdirectories.

For this assignment, your **HW06** folder should contain **task1.cu**, **task2.cu**, **matmul.cu**, and **stencil.cu**.

All commands or code must work on *Euler* with only the **nvidia/cuda** module loaded.

Since various commands may behave differently on your computer, we recommend that you test on *Euler* before you submit your homework.

Please submit clean code. Consider using a formatter like **clang-format**.

IMPORTANT: Before you begin, copy any provided files from **Assignments/HW06** directory of the **ECE 759 Resource Repo**. Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

github.com/kagawa16/epo759/tree/main/HW06

The GitHub link to your code folder should be: <https://github.com/YourGitHubName/epo759/HW06>

1. (a) Implement in a file called **matmul.cu** the **matmul** and **matmul_kernel** functions as declared and described in the comment section of **matmul.cuh**. These functions should compute the product of square matrices.
- (b) Write a program **task1.cu** which will complete the following (some memory management steps are omitted for clarity, but you should implement them in your code for it to work properly):
 - Create matrices (as 1D row major arrays) **A** and **B** of size **n** × **n** on the host.
 - Fill these matrices with random numbers in the range [-1, 1].
 - Prepare arrays that are allocated as device memory (they will be passed to your **matmul** function.)
 - Call your **matmul** function.
 - Print the last element of the resulting matrix.
 - Print the time taken to execute your **matmul** function in *milliseconds* using CUDA events.
 - Compile: **nvcc task1.cu matmul.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task1**
 - The executable is to be run as: **./task1 n threads_per_block**
 - The values **n** and **threads_per_block** are positive integers
 - The matrix size **n** is not necessarily a power of 2. We will not test with sizes larger than 2^{14} . For the number of threads, we will use multiples of 32.
 - Use **Slurm** to run your job on Euler
 - Example expected output:
-16.35
1.23
- (c) On an Euler *compute node*, run **task1** for each value **n** = $2^5, 2^6, \dots, 2^{14}$ and generate a plot **task1.pdf** that reports the time taken by your algorithm as a function of **n** when **threads_per_block** = 1024. Overlay another plot which plots the same relationship with a different **threads_per_block** of your choice.
- (d) **Going beyond the call of duty, do if you wish to:** Compare the scaling results with the results obtained in a previous assignment where you did a similar scaling analysis using a sequential implementation on the CPU. If you do this, drop your plots on Piazza. Please comment on the noted results.

2. (a) Implement in a file called `stencil.cu` the `stencil` and `stencil_kernel` functions as declared and described in the comment section of `stencil.cuh`. These functions should produce the 1D convolution of `image` and `mask` as the following:

$$\text{output}[i] = \sum_{j=-R}^R \text{image}[i+j] * \text{mask}[j+R] \quad i = 0, \dots, n-1.$$

Assume that `image[i] = 1` when $i < 0$ or $i > n-1$. Please read the comments in the `stencil.cuh` file to see which variables should be stored in shared memory. The use of shared memory is an important component of this problem.

- (b) Write a program `task2.cu` which will complete the following (some memory management steps are omitted for clarity, but you should implement them in your code):

- Create arrays `image` (length `n`), `output` (length `n`), and `mask` (length `2 * R + 1`) on the host.
- Fill the `image` and `mask` array with random numbers in the range $[-1, 1]$.
- Prepare arrays that are allocated as device memory (they will be passed to your `stencil` function.)
- Call your `stencil` function.
- Print the last element of the resulting `output` array.
- Print the time taken to execute your `stencil` function in *milliseconds* using CUDA events.
- Compile: `nvcc task2.cu stencil.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task2`
- Run via `Slurm` (where `n`, `R`, and `threads_per_block` are positive integers):
`./task2 n R threads_per_block`
- Example expected output:
11.36
1.23

- (c) On an Euler *compute node*, run `task2` for each value $n = 2^{10}, 2^{11}, \dots, 2^{29}$ and generate a plot `task2.pdf` which plots the time taken by your algorithm as a function of `n` when `threads_per_block` = 1024 and `R` = 128. Overlay another plot which plots the same relationship with a different `threads_per_block` of your choice.
- (d) *Going beyond the call of duty, do if you wish to:* Compare the scaling results with the results obtained in a previous assignment where you did a similar scaling analysis using a sequential implementation on the CPU. If you do this, drop your plots on Piazza. Please comment on the noted results.